

# Big Data Project

## Team members

- Ranim Mostafa Hamed – 22011576
- Rowan Fayed Mohamed – 22011454
- Sara Sameh Ahmed – 22011968
- Ola Ragab Saad Ali – 22010367
- Shorouq Tareq Ahmed – 22010350
- Nourhan Essam El-Din Mohamed – 22010287

## Data Lake Setup Summary

### 1. Data Lake Concept

A **Data Lake** is a centralized repository that allows you to store all structured and unstructured data at any scale. It supports multiple stages of data processing:

- **Bronze Layer:** Raw, unprocessed data (CSV files).
- **Silver Layer:** Cleaned and pre-processed data (Parquet files).
- **Gold Layer:** Final aggregated or analyzed results (e.g., factor analysis outputs).

### 2. HDFS Overview

**Hadoop Distributed File System (HDFS)** provides distributed storage, fault tolerance, and high-throughput access. It is used in this project to store cleaned and merged data in a structured way, allowing scalable processing.

### 3. MinIO Role

**MinIO** is used as an S3-compatible object storage to emulate a cloud-like data lake locally. Its benefits in the project include:


- Secure storage for raw (Bronze) and processed (Silver/Gold) data.

- Easy integration with Python for automated uploads.
- Acts as an intermediary before moving cleaned data into HDFS.

## 4. Docker Setup

**Docker** containers are used to encapsulate all services, ensuring reproducibility and isolation. The project uses:

- **MinIO Container:** Hosts Bronze, Silver, and Gold buckets.



A screenshot of a terminal window displaying a Docker Compose configuration file. The file defines a service named 'minio' which runs the 'minio/minio:latest' image. The container is named 'minio-server'. It exposes ports 9002 and 9001. Environment variables 'MINIO\_ROOT\_USER' and 'MINIO\_ROOT\_PASSWORD' are set to 'minioadmin' and 'minioadmin123' respectively. A volume 'minio-data:/data' is mounted. The command for the container is 'server /data --console-address ":9001"'. A healthcheck is defined with a curl test to 'http://localhost:9000/minio/health/live' every 30 seconds with a 20-second timeout and 3 retries. The container is connected to a network named 'bigdata-network'.

```
1  services:
2    # MinIO Object Storage
3    minio:
4      image: minio/minio:latest
5      container_name: minio-server
6      ports:
7        - "9002:9000"
8        - "9001:9001"
9      environment:
10        MINIO_ROOT_USER: minioadmin
11        MINIO_ROOT_PASSWORD: minioadmin123
12      volumes:
13        - minio-data:/data
14      command: server /data --console-address ":9001"
15      healthcheck:
16        test: ["CMD", "curl", "-f", "http://localhost:9000/minio/health/live"]
17        interval: 30s
18        timeout: 20s
19        retries: 3
20      networks:
21        - bigdata-network
```

- **HDFS NameNode Container:** Manages HDFS metadata.

```
● ● ●

1 # HDFS NameNode
2 namenode:
3   image: bde2020/hadoop-namenode:2.0.0-hadoop3.2.1-java8
4   container_name: hdfs-namenode
5   ports:
6     - "9870:9870"      # NameNode Web UI
7     - "9000:9000"      # HDFS API (exposed on different internal port)
8   environment:
9     - CLUSTER_NAME=bigdata-cluster
10    - CORE_CONF_fs_defaultFS=hdfs://namenode:9000
11    - CORE_CONF_hadoop_http_staticuser_user=root
12    - HDFS_CONF_dfs_replication=1
13    - HDFS_CONF_dfs_namenode_datanode_registration_ip__hostname__check=false
14   volumes:
15     -.hadoop-namenode:/hadoop/dfs/name
16   networks:
17     - bigdata-network
18   healthcheck:
19     test: ["CMD", "curl", "-f", "http://localhost:9870"]
20     interval: 30s
21     timeout: 10s
22     retries: 3
```

- **HDFS DataNode Container:** Stores actual data blocks.

```
● ● ●

1 # HDFS DataNode
2 datanode:
3   image: bde2020/hadoop-datanode:2.0.0-hadoop3.2.1-java8
4   container_name: hdfs-datanode
5   ports:
6     - "9864:9864"      # DataNode Web UI
7   environment:
8     - CORE_CONF_fs_defaultFS=hdfs://namenode:9000
9     - CORE_CONF_hadoop_http_staticuser_user=root
10    - HDFS_CONF_dfs_replication=1
11   volumes:
12     -.hadoop-datanode:/hadoop/dfs/data
13   networks:
14     - bigdata-network
15   depends_on:
16     - namenode
```

## Benefits of Docker in the project:

- Simplifies setup of complex distributed systems.

- Ensures consistent environment across machines.
- Allows running HDFS and MinIO locally without manual installations.

## 5. Data Flow and Upload Process

### 1. Raw data upload (Bronze):

CSV files (`london_weather_dataset2.csv`, `london_traffic_dataset2.csv`) are uploaded to the MinIO Bronze bucket using Python scripts with the MinIO SDK.

### 2. Cleaned data upload (Silver):

Processed Parquet files (`weather_cleaned.parquet`, `traffic_cleaned.parquet`, `merged_data(3).parquet`) are uploaded to the MinIO Silver bucket.


### 3. Transfer to HDFS:

Data is copied from Silver to HDFS in three directories:

- `/bigdata/weather`
- `/bigdata/traffic`
- `/bigdata/merged`

### 4. Method used:

Python automates Docker commands to execute native Hadoop CLI (`hdfs dfs -put`) inside the HDFS NameNode container. This avoids WebHDFS dependency and ensures reliable ingestion in the containerized environment.



```

1 # Create HDFS directories first
2 print("Creating HDFS directory structure...")
3 directories = ["/bigdata", "/bigdata/weather", "/bigdata/traffic", "/bigdata/merged"]
4
5 for directory in directories:
6     mkdir_cmd = f'docker exec hdfs-namenode hdfs dfs -mkdir -p {directory}'
7     result = subprocess.run(mkdir_cmd, shell=True, capture_output=True, text=True)
8     if result.returncode == 0:
9         print(f"Created/verified: {directory}")
10    else:
11        if "File exists" not in result.stderr:
12            print(f"Warning for {directory}: {result.stderr.strip()}")
13

```



```
1 try:
2     # Step 1: Copy file into container
3     temp_name = local_file.replace(" ", "_").replace("(, )", "").replace(")", "")
4     copy_cmd = f'docker cp "{local_path}" hdfs-namenode:/tmp/{temp_name}'
5     result = subprocess.run(copy_cmd, shell=True, capture_output=True, text=True)
6
7     if result.returncode != 0:
8         print(f" Failed to copy to container: {result.stderr}")
9         continue
10
11    # Step 2: Put file into HDFS
12    put_cmd = f'docker exec hdfs-namenode hdfs dfs -put -f /tmp/{temp_name} {hdfs_path}'
13    result = subprocess.run(put_cmd, shell=True, capture_output=True, text=True)
14
15    if result.returncode != 0:
16        print(f" Failed to upload to HDFS: {result.stderr}")
17        continue
18
19    # Step 3: Clean up temp file
20    clean_cmd = f'docker exec hdfs-namenode rm /tmp/{temp_name}'
21    subprocess.run(clean_cmd, shell=True, capture_output=True, text=True)
22
23    print(f" Uploaded successfully")
24    success_count += 1
25
26 except Exception as e:
27     print(f"Error: {e}")
28
```

## 5. Results upload (Gold):

Analysis outputs (CSV and DOCX) are uploaded to MinIO Gold bucket for further use or reporting.

## 6. Python Automation

Python is used as an **orchestration layer** to:

- Upload files to MinIO buckets.
- Automate creation of HDFS directories and file uploads.
- Verify bucket contents and HDFS structure.
- Scripted in `complete_setup.py` for the full pipeline.

The **actual data movement in HDFS** is performed by native `hdfs dfs` commands.

## 7. Security Considerations

### 1. MinIO Security:

- **Authentication:** MinIO access is protected using an access key (`MINIO_ACCESS_KEY`) and a secret key (`MINIO_SECRET_KEY`). Only

authorized users can upload or read data from Bronze, Silver, or Gold buckets.

- **Bucket Isolation:** Data is organized in separate buckets (`bronze`, `silver`, `gold`) to prevent accidental data mix-up and allow fine-grained access control.
- **HTTPS (optional):** MinIO supports TLS/SSL for secure communication (though in the local Docker setup we used HTTP for simplicity).

## 2. HDFS Security:

- **User Isolation:** All HDFS operations are executed under the user `root` to maintain control over file ownership and permissions.
- **Directory Structure:** Cleaned data is stored in dedicated HDFS directories (`/bigdata/weather`, `/bigdata/traffic`, `/bigdata/merged`), reducing risk of overwriting unrelated data.
- **Overwrite Control:** The `hdfs dfs -put -f` command ensures only intended files are overwritten.

## 3. Docker Security:

- **Container Isolation:** Each service runs in its own container (`minio-server`, `hdfs-namenode`, `hdfs-datanode`), isolating processes and reducing cross-service interference.
- **Network Segmentation:** All containers communicate via a dedicated Docker network (`bigdata-network`), avoiding exposure to other host services.

## 4. Data Integrity:

- **Checks during upload:** Scripts verify the existence of files before upload and confirm successful transfer.
- **Verification phase:** The `verify_setup()` script ensures all buckets and HDFS directories contain the expected files, helping detect corruption or misplacement early.

## 5. Benefits Realized:

- Secure, authenticated access to sensitive datasets.
- Clear separation between raw, processed, and result data layers.

- Safe and reproducible environment thanks to Docker containerization.

## 8. Project Architecture

Local Data Directory



MinIO (Bronze/Silver/Gold)



HDFS NameNode / DataNode



Big Data Analysis & Gold Layer

### Key points:

- Bronze: Raw CSV

Name	Last Modified
london_traffic_dataset2.csv	Today, 03:28
london_weather_dataset2.csv	Today, 03:28

- Silver: Cleaned Parquet

Name	Last Modified
merged_data (3).parquet	Today, 03:28
traffic_cleaned.parquet	Today, 03:28
weather_cleaned.parquet	Today, 03:28

- HDFS: Distributed storage for Silver data

The screenshot shows a web-based HDFS browser interface. The title bar includes tabs for "BigData\_FinalProject\_LU[1].pdf", "MinIO Console", and "Browsing HDFS". The main navigation bar has links for "Hadoop", "Overview", "Datanodes", "Datanode Volume Failures", "Snapshot", "Startup Progress", and "Utilities". Below the navigation is a search bar and a "Browse Directory" section. The directory listing table has columns for "Permission", "Owner", "Group", "Size", "Last Modified", "Replication", "Block Size", and "Name". A single entry is listed: "drwxr-xr-x" permission, owner "root", group "supergroup", size 0 B, last modified Dec 14 03:28, replication 0, block size 0 B, and name "bigdata". The footer displays "Hadoop, 2019."



Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	root	supergroup	0 B	Dec 14 03:28	0	0 B	bigdata


This screenshot shows the contents of the "/bigdata" directory. The title bar shows the URL "localhost:9870/explorer.html#/bigdata". The main content area is titled "Browse Directory" with the path "/bigdata" in the address bar. The table lists three entries: "merged", "traffic", and "weather", all with the same permissions and details as the root directory.

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	root	supergroup	0 B	Dec 14 03:29	0	0 B	merged
drwxr-xr-x	root	supergroup	0 B	Dec 14 03:29	0	0 B	traffic
drwxr-xr-x	root	supergroup	0 B	Dec 14 03:29	0	0 B	weather


This screenshot shows the contents of the "/bigdata/weather" directory. The title bar shows the URL "localhost:9870/explorer.html#/bigdata/weather". The main content area is titled "Browse Directory" with the path "/bigdata/weather" in the address bar. It lists a single file, "weather\_cleaned.parquet", which is 186.36 KB in size and was last modified on Dec 14 03:29. The replication factor is 1 and the block size is 128 MB.

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	root	supergroup	186.36 KB	Dec 14 03:29	1	128 MB	weather_cleaned.parquet





- Gold: Analysis results



- Python: Automation & orchestration
- Docker: Encapsulation & reproducibility
- MinIO: Local object storage simulation

## 9. Summary

This setup demonstrates a **complete Data Lake workflow**:

- Data ingestion from raw CSV to object storage (MinIO).
- Data cleaning and storage in Parquet format.

- Reliable ingestion into HDFS for scalable processing.
- Automated verification to ensure all components function correctly.

**Outcome:** All datasets are securely stored, organized, and ready for further big data analysis, with Docker ensuring the project is fully reproducible and portable.

# Data Preprocessing

## Inspect Initial Dataset

```
def analyze_dataset_before_cleaning(df, dataset_name):
    """Analyze dataset before cleaning"""
    print(f"ANALYZING {dataset_name.upper()} BEFORE CLEANING")
    print('*'*60)

    print(f"Total rows: {len(df)}")
    print(f"Duplicate rows: {df.duplicated().sum()}")
    print(f"Total missing values: {df.isnull().sum().sum()}")

    # Analyze each column
    for column in df.columns:
        print(f"\n{column}:")
        print(f"  Type: {df[column].dtype}")
        print(f"  Missing: {df[column].isnull().sum()} ({(df[column].isnull().sum() / len(df)) * 100:.1f}%)"

        if df[column].dtype in ['int64', 'float64']:
            print(f"  Min: {df[column].min():.2f}")
            print(f"  Max: {df[column].max():.2f}")
            print(f"  Mean: {df[column].mean():.2f}")
            print(f"  Std: {df[column].std():.2f}")

        if df[column].dtype == 'object':
            unique_count = df[column].nunique()
            print(f"  Unique values: {unique_count}")
            if unique_count < 20: # Show unique values for small sets
                print(f"  Values: {df[column].dropna().unique()[:10]}")

    return df
```

## 1. Weather Dataset Analysis & Cleaning

### 1.1 Initial Data Quality

**Dataset:**

- Total rows: 5,000
- Duplicate rows: 0
- Total missing values: 973 (across all columns)

## Column Analysis:

Column	Data Type	Missing	Range/Values	Issues Identified
weather_id	float64	47 (0.9%)	5001-9900	Nulls, duplicates
date_time	object	53 (1.1%)	4,886 unique	Mixed formats
season	object	50 (1.0%)	Winter/Spring/Summer	Nulls
city	object	59 (1.2%)	London	Nulls
temperature_c	float64	104 (2.1%)	-30°C to 60°C	Extreme outliers
humidity	float64	85 (1.7%)	-10% to 150%	Invalid range
rain_mm	float64	85 (1.7%)	0-120mm	Extreme values
wind_speed_kmh	float64	92 (1.8%)	0.16-250 km/h	Unrealistic peaks
visibility_m_weather	object	104 (2.1%)	Mixed	String contamination
weather_condition	object	48 (1.0%)	5 categories	Nulls
air_pressure_hpa	float64	246 (4.9%)	978-1045 hPa	Highest missing rate

## 1.2 Cleaning

### Step 1: Date Format Standardization

```
# 1. Fix date format
print("\n1. FIXING DATE FORMATS")
df['date_time'] = pd.to_datetime(df['date_time'], errors='coerce', format='%Y-%m-%d %H:%M:%S')
```

- Converted all dates to datetime format

## Step 2: Null Date Removal

```
# 2. check null dates
print("\n2. Remove Null DATES")
null_dates = df['date_time'].isna().sum()
print(f"Null dates before dropping: {null_dates}")

df = df.dropna(subset=['date_time']).reset_index(drop=True)
null_dates_after = df['date_time'].isna().sum()
print(f"Null dates after dropping: {null_dates_after}")
```

- Null dates before: 250
- Null dates after: 0

## Step 3: Duplicate Removal

```
# 3. Remove duplicates
print("\n3. REMOVING DUPLICATES")
duplicates = df.duplicated().sum()
df = df.drop_duplicates()
print(f" Removed {duplicates} duplicate rows")
```

- No complete duplicate rows found (0 removed)

## Step 4: Temperature Outlier Correction

```
# 4. Fix temperature outliers
print("\n4. FIXING TEMPERATURE OUTLIERS")
temp_outliers = df[(df['temperature_c'] < -15) | (df['temperature_c'] > 40)].shape[0]
temp_min_before = df['temperature_c'].min()
temp_max_before = df['temperature_c'].max()

df['temperature_c'] = df['temperature_c'].clip(-15, 40)

temp_min_after = df['temperature_c'].min()
temp_max_after = df['temperature_c'].max()

print(f" Fixed {temp_outliers} temperature outliers")
print(f" Range before: {temp_min_before:.1f} to {temp_max_before:.1f}°C")
print(f" Range after: {temp_min_after:.1f} to {temp_max_after:.1f}°C")
```

- Applied realistic bounds for London climate: -15°C to 40°C
- Fixed 95 temperature outliers
- Range before: -30.0°C to 60.0°C
- Range after: -15.0°C to 40.0°C

## Step 5: Humidity Correction

```
# 5. Fix humidity (0-100%)
print("\n5. FIXING HUMIDITY VALUES")
hum_outliers = df[(df['humidity'] < 0) | (df['humidity'] > 100)].shape[0]
hum_min_before = df['humidity'].min()
hum_max_before = df['humidity'].max()

df['humidity'] = df['humidity'].clip(0, 100)

hum_min_after = df['humidity'].min()
hum_max_after = df['humidity'].max()

print(f" Fixed {hum_outliers} humidity outliers")
print(f" Range before: {hum_min_before:.1f} to {hum_max_before:.1f}%")
print(f" Range after: {hum_min_after:.1f} to {hum_max_after:.1f}%")
```

- Clipped values to valid percentage range: 0-100%
- Fixed 94 humidity outliers
- Range before: -10.0% to 150.0%
- Range after: 0.0% to 100.0%

## Step 6: Missing Value Imputation

```
# 6. Handle missing values
print("\n6. HANDLING MISSING VALUES")
# Keep only rows where weather_id is NOT null
df = df[df['weather_id'].notna()]

# Replace 'Unknown' with NaN first
columns_to_process = [col for col in df.columns if col != 'weather_id']
df[columns_to_process] = df[columns_to_process].replace('Unknown', pd.NA)

# Then impute missing values
for col in columns_to_process:
    missing_before = df[col].isnull().sum()
    if missing_before > 0:
        df[col] = df[col].fillna(df[col].mode().iloc[0])

    missing_after = df[col].isnull().sum()
    print(f" {col}: {missing_before} missing + {missing_after} missing")
```

- Removed rows where weather\_id was null (primary key)
- Replaced "Unknown" strings with NaN
- Applied mode imputation for remaining missing values

Column	Missing Before	Missing After	Imputation Method
--------	----------------	---------------	-------------------

season	47	0	Mode
city	54	0	Set to "London"
temperature_c	97	0	Mode
humidity	82	0	Mode
rain_mm	81	0	Mode
wind_speed_kmh	89	0	Mode
visibility_m_weather	146	0	Mode
weather_condition	44	0	Mode
air_pressure_hpa	236	0	Mode

## Step 7: Data Type Standardization

```
# 7. Handle Data types
print("\n7. FINAL DATA TYPES")

# Ensure 'city' is properly handled - set to 'London' for all rows
df['city'] = 'London'
df['temperature_c'] = df['temperature_c'].round(1)
df['humidity'] = df['humidity'].round().astype(int)
df['rain_mm'] = df['rain_mm'].round(1)
df['visibility_m_weather'] = df['visibility_m_weather'].astype(float)

# Display final data types
print(df.dtypes)
```

- Rounded temperature to 1 decimal place
- Converted humidity to integer
- Rounded rain to 1 decimal place
- Converted visibility to float

## Step 8: Final Visibility Correction

```
# 8. Fix Visibility values  
df['visibility_m_weather'] = df['visibility_m_weather'].clip(lower=0)
```

- Applied lower bound of 0 meters (no negative visibility)

## 1.3 Weather Cleaning Results

### Final Statistics:

- Original rows: 5,000
- Final rows: 4,704
- Rows removed: 296
- Data retained: **94.1%**

### Final Data Types:

```
weather_id                  float64  
date_time                   datetime64[ns]  
season                       object  
city                          object  
temperature_c                float64  
humidity                     int64  
rain_mm                      float64  
wind_speed_kmh               float64  
visibility_m_weather          float64  
weather_condition            object  
air_pressure_hpa              float64  
dtype: object
```

## 2. Traffic Dataset Analysis & Cleaning

### 2.1 Initial Data Quality

#### Dataset:

- Total rows: 5,000
- Duplicate rows: 0
- Total missing values: 693 (across all columns)

#### Column Analysis:

Column	Data Type	Missing	Range/Values	Issues Identified
traffic_id	float64	48 (1.0%)	9001-13900	Nulls, duplicates
date_time	object	56 (1.1%)	4,874 unique	Mixed formats
city	object	51 (1.0%)	London	Nulls
area	object	48 (1.0%)	7 districts	Nulls
vehicle_count	float64	108 (2.2%)	45.9-25,000	Extreme outliers
avg_speed_kmh	float64	81 (1.6%)	-50 to 64.85	Negative values
accident_count	float64	46 (0.9%)	0-55	Unrealistic peaks
congestion_level	object	45 (0.9%)	High/Medium/Low	Nulls
road_condition	object	112 (2.2%)	Dry/Wet/Snowy	Highest missing
visibility_m_traffic	float64	98 (2.0%)	-34 to 10,099m	Negative values

## 2.2 Cleaning

### Step 1: Date Format Standardization

```
# 1. Fix date format
print("\n1. FIXING DATE FORMATS")
df['date_time'] = pd.to_datetime(df['date_time'], errors='coerce', format='%Y-%m-%d %H:%M:%S')
```

- Same process as weather dataset
- Result: 250 invalid dates identified

### Step 2: Null Date Removal

```

# 2. check null dates
print("\n2. Remove Null DATES")
null_dates = df['date_time'].isna().sum()
print(f"Null dates before dropping: {null_dates}")

df = df.dropna(subset=['date_time']).reset_index(drop=True)
null_dates_after = df['date_time'].isna().sum()
print(f"Null dates after dropping: {null_dates_after}")

```

- Null dates before: 250
- Null dates after: 0

### Step 3: Duplicate Removal

```

# 3. Remove duplicates
print("\n3. REMOVING DUPLICATES")
duplicates = df.duplicated().sum()
df = df.drop_duplicates()
print(f" Removed {duplicates} duplicate rows")

```

- No complete duplicate rows found (0 removed)

### Step 4: Speed Outlier Removal (IQR Method)

```

# 4. Fix speed outliers
print("\n4. FIXING SPEED VALUES")

len_before = len(df)
Q1 = df['avg_speed_kmh'].quantile(0.25)
Q3 = df['avg_speed_kmh'].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filter
df = df[
    (df['avg_speed_kmh'] >= lower_bound) &
    (df['avg_speed_kmh'] <= upper_bound)
]

print("Removed rows:", len_before - len(df))

```

- Applied Interquartile Range (IQR) method for outlier detection
- Calculated Q1 (25th percentile) and Q3 (75th percentile)
- Defined bounds:  $Q1 - 1.5 \times IQR$  to  $Q3 + 1.5 \times IQR$
- **Removed 128 rows** with unrealistic speeds (including negative values)

### Step 5: Vehicle Count Outlier Removal (IQR Method)

```

# 5. Fix vehicle count
print("\n5. FIXING VEHICLE COUNT")

len_before = len(df)
Q1 = df['vehicle_count'].quantile(0.25)
Q3 = df['vehicle_count'].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filter
df = df[
    (df['vehicle_count'] >= lower_bound) &
    (df['vehicle_count'] <= upper_bound)
]

print("Removed rows:", len_before - len(df))

```

- Applied IQR
- **Removed 162 rows** with extreme vehicle counts

## Step 6: Missing Value Imputation

```

# 6. Handle missing values
print("\n7. HANDLING MISSING VALUES")
# Keep only rows where weather_id is NOT null
df = df[df['traffic_id'].notna()]

# Replace 'Unknown' with NaN first
columns_to_process = [col for col in df.columns if col != 'traffic_id']
df[columns_to_process] = df[columns_to_process].replace('Unknown', pd.NA)

# Then impute missing values
for col in columns_to_process:
    missing_before = df[col].isnull().sum()
    if missing_before > 0:
        df[col] = df[col].fillna(df[col].mode().iloc[0])

    missing_after = df[col].isnull().sum()
    print(f" {col}: {missing_before} missing → {missing_after} missing")

```

- Kept only rows where traffic\_id was not null
- Applied mode imputation for categorical and numerical fields

Column	Missing Before	Missing After	Imputation Method
city	40	0	Set to "London"

area	43	0	Mode
accident_count	43	0	Mode
congestion_level	41	0	Mode
road_condition	101	0	Mode
visibility_m_traffic	85	0	Mode

## Step 7: Data Type Standardization

```
# 7. Handle Data types
print("\n8. FINAL DATA TYPES")
df['city'] = 'London'
df['vehicle_count'] = df['vehicle_count'].astype(int)
df['accident_count'] = df['accident_count'].astype(int)
df['avg_speed_kmh'] = df['avg_speed_kmh'].round(1)
df['visibility_m_traffic'] = df['visibility_m_traffic'].astype(float)

# Display final data types
print(df.dtypes)
```

- Ensured city is "London" for all rows
- Converted vehicle\_count to integer
- Converted accident\_count to integer
- Rounded avg\_speed\_kmh to 1 decimal
- Converted visibility to float

## Step 8: Visibility Correction

```
# 8. Fix Visibility values
df['visibility_m_traffic'] = df['visibility_m_traffic'].clip(lower=0)
```

- Clipped to minimum of 0 meters

## Step 9: Speed Correction

```
# 9. Fix Speed values  
df['avg_speed_kmh'] = df['avg_speed_kmh'].clip(lower=0)
```

- Clipped to minimum of 0

## 2.3 Traffic Cleaning Results

### Final Statistics:

- Original rows: 5,000
- Final rows: 4,420
- Rows removed: 580
- Data retained: **88.4%**

### Final Data Types:

```
traffic_id                      float64  
date_time                       datetime64[ns]  
city                            object  
area                            object  
vehicle_count                   int64  
avg_speed_kmh                  float64  
accident_count                 int64  
congestion_level               object  
road_condition                 object  
visibility_m_traffic           float64  
dtype: object
```

## 3. Key Findings & Observations

### 3.1 Data Loss Analysis

#### Weather Dataset:

- 5.9% data loss (296 rows)
- Primary causes: Invalid dates, null primary keys.

#### Traffic Dataset:

- 11.6% data loss (580 rows)
- Primary causes: Invalid dates, speed outliers and vehicle count outliers .

### 3.2 Data Quality Improvements

#### Before Cleaning:

- Mixed date formats
- Invalid ranges
- Missing values
- Outliers skewing distributions and means

#### After Cleaning:

- Standardized datetime format
- Realistic value ranges for all measurements
- Complete records (no missing values)

## 4. Output Files Generated

```
: weather_cleaned.to_parquet('weather_cleaned.parquet', index=False)
traffic_cleaned.to_parquet('traffic_cleaned.parquet', index=False)
```

1. weather\_cleaned.parquet - Cleaned weather data (4,704 rows)
2. traffic\_cleaned.parquet - Cleaned traffic data (4,420 rows)

## Datasets Merging :

1- Defining The function `check_merge_ready` that performs a quick check on a DataFrame before merging :

1. Check that key columns (`key_columns`) exist in the DataFrame.
2. Check for missing values in the key columns.
3. Check for duplicate rows in the DataFrame.
4. Verify the `date_time` column type (if present) is datetime.
5. Standardize the `city` column (convert to string, strip spaces, title case) and show unique values.

The main goal is to ensure the dataset is merge-ready without issues in columns, values, or formatting.

### Data cleaning validation before merging

```
def check_merge_ready(df, dataset_name, key_columns):
    print(f"\nChecking {dataset_name} for merge readiness...")
    print("-" * 50)

    # 1. Check key columns exist
    for col in key_columns:
        if col not in df.columns:
            print(f"Missing key column: {col}")
        else:
            print(f"Column '{col}' exists")

    # 2. Check for missing values in key columns
    for col in key_columns:
        missing = df[col].isnull().sum()
        if missing > 0:
            print(f"Column '{col}' has {missing} missing values")
        else:
            print(f"Column '{col}' has no missing values")

    # 3. Check duplicates
    dup_count = df.duplicated().sum()
    if dup_count > 0:
        print(f"Dataset has {dup_count} duplicate rows")
    else:
        print("No duplicate rows")

    # 4. Check date_time type if in keys
    if 'date_time' in key_columns:
        if not pd.api.types.is_datetime64_any_dtype(df['date_time']):
            print("'date_time' column is NOT datetime type")
        else:
            print("'date_time' column is datetime type")

    # 5. Standardize city column if in keys
    if 'city' in key_columns:
        df['city'] = df['city'].astype(str).str.strip().str.title()
        unique_cities = df['city'].unique()
        print(f"'city' standardized. Unique cities: {unique_cities}")
```

After running it we found :

- Both `date_time` and `city` columns exist in each dataset.
- There are no missing values in the key columns.
- There are no duplicate rows.
- The `date_time` column is confirmed to be datetime type.
- The `city` column has been standardized (converted to title case), and both datasets only contain ['London'] as the unique city.

```

# Define merge keys
merge_keys = ['date_time', 'city']

# Run checks
check_merge_ready(weather_cleaned, "Weather Dataset", merge_keys)
check_merge_ready(traffic_cleaned, "Traffic Dataset", merge_keys)

Checking Weather Dataset for merge readiness...
-----
Column 'date_time' exists
Column 'city' exists
Column 'date_time' has no missing values
Column 'city' has no missing values
No duplicate rows
'date_time' column is datetime type
'city' standardized. Unique cities: ['London']

Checking Traffic Dataset for merge readiness...
-----
Column 'date_time' exists
Column 'city' exists
Column 'date_time' has no missing values
Column 'city' has no missing values
No duplicate rows
'date_time' column is datetime type
'city' standardized. Unique cities: ['London']

```

We converted the **date\_time** columns in both the **weather\_cleaned** and **traffic\_cleaned** DataFrames into datetime objects using `pd.to_datetime()`. Then, they round down each timestamp to the nearest hour using `.dt.floor('H')`. This ensures that all timestamps are aligned at the hour level, which is useful for merging.

```

weather_cleaned['date_time'] = pd.to_datetime(weather_cleaned['date_time']).dt.floor('H')
traffic_cleaned['date_time'] = pd.to_datetime(traffic_cleaned['date_time']).dt.floor('H')

```

Finally we merged the 2 datasets on “date\_time” and “city” columns

```

merged_df = pd.merge(
    traffic_cleaned,
    weather_cleaned,
    on=['city', 'date_time'],
    how='inner'
)
merged_df

```

	traffic_id	date_time	city	area	vehicle_count	avg_speed_kmh	accident_count	congestion_level	road_condition	visibility_m_traffic	weather_id	season
0	9002.0	2020-01-01 01:00:00	London	Westminster	267	43.9	0	Medium	Dry	9947.0	5002.0	Winter
1	9003.0	2020-01-01 02:00:00	London	Camden	188	59.4	0	Low	Dry	10072.0	5003.0	Winter
2	9004.0	2020-01-01 04:00:00	London	Greenwich	226	47.4	0	Low	Dry	10047.0	5005.0	Winter
3	9005.0	2020-01-01 04:00:00	London	Southwark	340	47.3	0	Low	Dry	10080.0	5005.0	Winter
4	9006.0	2020-01-01 06:00:00	London	Westminster	350	39.8	0	Medium	Wet	5534.0	5007.0	Winter

# Factor Analysis on Urban Traffic and Weather Data

The goal of this project is to analyze urban traffic data in London under varying weather conditions, using factor analysis to reduce dimensionality and identify latent factors that influence traffic behavior. The dataset contains 4162 records with 15 initial features covering traffic volume, road and weather conditions, and temporal information.

## 1-Data Cleaning and Preprocessing

### Initial Cleaning

- *Removed non-informative columns: city, date\_time, traffic\_id, weather\_id, area.*

```
1 # 1. Drop columns that are constant or highly unique IDs/Text and won't add variance to the FA
2 # 'city' is constant (London)
3 # 'date_time' is temporal (excluding for FA)
4 # 'traffic_id', 'weather_id' are IDs
5 # 'area' is categorical (Nominal) - excluded initially to simplify FA
6 df_cleaned = df.drop(columns=['city', 'date_time', 'traffic_id', 'weather_id', 'area'])
```

- 
- *Checked for nulls; dataset contains no missing values after cleaning.*

Columns after cleaning (15 features):

```
[30] ✓ 0s df_cleaned.columns
...
Index(['vehicle_count', 'avg_speed_kmh', 'accident_count',
       'visibility_m_traffic', 'temperature_c', 'humidity', 'rain_mm',
       'wind_speed_kmh', 'visibility_m_weather', 'air_pressure_hpa', 'hour',
       'congestion_level_encoded', 'season_encoded', 'road_severity',
       'weather_severity'],
      dtype='object')
```

## 2- Encoding Categorical Variables

- `congestion_level` → ordinal: `Low=1, Medium=2, High=3`
- `season` → ordinal: `Winter=1, Spring=2, Summer=3, Autumn=4`
- `road_condition` → ordinal: `Dry=0, Wet=1, Snowy=2`
- `weather_condition` → ordinal: `Clear=0, Partly Cloudy=1, Fog=2, Rain=3, Snow=4, Storm=5`

```
1 # --- Handling Categorical Variables ---
2
3 # 1. Ordinal Encoding for 'congestion_level' (Label Encoding)
4 congestion_map = {'Low': 1, 'Medium': 2, 'High': 3}
5 df_cleaned['congestion_level_encoded'] = df_cleaned['congestion_level'].map(congestion_map)
6 df_cleaned.drop(columns=['congestion_level'], inplace=True)
7
8 # 2. Ordinal Encoding for 'season' (Label Encoding)
9 season_map = {'Winter': 1, 'Spring': 2, 'Summer': 3, 'Autumn': 4}
10 df_cleaned['season_encoded'] = df_cleaned['season'].map(season_map)
11 df_cleaned.drop(columns=['season'], inplace=True)
12
13 road_severity_map = {
14     'Dry': 0,
15     'Wet': 1,
16     'Snowy': 2
17 }
18 df_cleaned['road_severity'] = df_cleaned['road_condition'].map(road_severity_map)
19
20 # Weather severity
21 weather_severity_map = {
22     'Clear': 0,
23     'Partly Cloudy': 1,
24     'Fog': 2,
25     'Rain': 3,
26     'Snow': 4,
27     'Storm': 5
28 }
29 df_cleaned['weather_severity'] = df_cleaned['weather_condition'].map(weather_severity_map)
```

All original categorical columns were dropped after encoding.

```
1 df_cleaned = df_cleaned.drop(columns=['road_condition', 'weather_condition'])
```

**Final dataset features (15 numeric columns):**

*Columns after Encoding:*

['vehicle\_count', 'avg\_speed\_kmh', 'accident\_count', 'visibility\_m\_traffic', 'temperature\_c', 'humidity',  
'rain\_mm', 'wind\_speed\_kmh', 'visibility\_m\_weather', 'air\_pressure\_hpa', 'hour',  
'congestion\_level\_encoded', 'season\_encoded', 'road\_severity', 'weather\_severity']

### 3 -Standardization

All features were standardized using StandardScaler to ensure uniform scaling for factor analysis.

```
● ● ●  
1 scaler = StandardScaler()  
2 df_scaled_array = scaler.fit_transform(df_cleaned)  
3 df_scaled = pd.DataFrame(df_scaled_array, columns=df_cleaned.columns)
```

Example of first 5 rows after scaling:

```
Data Standardization Complete. First 5 rows of scaled data:  
   vehicle_count  avg_speed_kmh  accident_count  visibility_m_traffic  \  
0      -1.058205      0.961524       -0.14341        0.745965  
1      -1.130305      2.268819       -0.14341        0.782496  
2      -1.095624      1.256720       -0.14341        0.775190  
3      -0.991581      1.248286       -0.14341        0.784834  
4      -0.982455      0.615724       -0.14341       -0.543729  
  
   temperature_c  humidity  rain_mm  wind_speed_kmh  visibility_m_weather  \  
0      -1.078642  -1.106672  -0.256140       0.001473        0.573310  
1      -0.812038  -0.637482  -0.256140      -0.274728        0.573310  
2      -1.232992  -0.972617  -0.256140      -0.010498        0.573310  
3      -1.232992  -0.972617  -0.256140      -0.010498        0.573310  
4      -0.419147    0.099816   0.013793      -0.407209       -1.231632
```

No missing values remained after scaling.

#### 4. Suitability for Factor Analysis


KMO and Bartlett's Test

- KMO Measure: **0.646** → acceptable sampling adequacy
- Bartlett's Test: **p-value < 0.001** → correlation matrix is not an identity

```
● ● ●  
1 kmo_all, kmo_model = calculate_kmo(df_scaled)  
2 bartlett_spher, bartlett_p = calculate_bartlett_sphericity(df_scaled)  
3  
4 print("\n--- Suitability Tests for Factor Analysis ---")  
5 print(f"Kaiser-Meyer-Olkin (KMO) Measure: {kmo_model:.3f}")  
6 print(f"Bartlett's Test P-value: {bartlett_p:.3e}")
```

## 5- Determining Number of Factors

- Eigenvalues >1 suggested **6 factors**, but based on interpretability, **4 factors were selected**.
- Scree plot confirmed an elbow at factor 4.



## 6. Factor Analysis (Varimax Rotation)

Performed factor analysis on the standardized dataset with 4 factors using Varimax rotation.

```
● ● ●

1 # --- Final Factor Analysis with Varimax Rotation ---
2
3 num_factors = 4
4 fa_final = FactorAnalyzer(n_factors=num_factors, rotation='varimax')
5 fa_final.fit(df_scaled)
6
7 # Factor Loadings DataFrame
8 factor_loadings = pd.DataFrame(
9     fa_final.loadings_,
10    index=df_scaled.columns,
11    columns=[f'Factor {i+1}' for i in range(num_factors)])
12 )
```

Factor Loadings (Threshold  $\geq 0.4$  Highlighted)

		Factor 1	Factor 2	Factor 3	Factor 4
vehicle_count		<b>0.923000</b>	-0.058000	0.228000	-0.036000
avg_speed_kmh		<b>-0.792000</b>	-0.231000	<b>0.566000</b>	0.034000
accident_count		0.001000	-0.005000	-0.065000	0.004000
visibility_m_traffic		-0.112000	-0.256000	<b>0.634000</b>	-0.012000
temperature_c		0.146000	-0.020000	0.070000	<b>0.986000</b>
humidity		0.016000	<b>0.462000</b>	-0.070000	0.033000
rain_mm		-0.001000	0.239000	-0.007000	0.002000
wind_speed_kmh		0.010000	0.081000	0.012000	-0.017000
visibility_m_weather		-0.037000	<b>-0.404000</b>	0.219000	-0.039000
air_pressure_hpa		-0.006000	-0.002000	-0.014000	-0.003000
hour		0.392000	0.004000	0.081000	0.141000
congestion_level_encoded		<b>0.708000</b>	0.162000	<b>-0.433000</b>	-0.028000
season_encoded		-0.001000	-0.004000	-0.011000	<b>0.405000</b>
road_severity		0.127000	0.376000	<b>-0.486000</b>	-0.028000
weather_severity		0.040000	<b>0.996000</b>	-0.062000	0.010000

#### Interpretation of Factors:

1. Factor 1 – Traffic Load & Congestion:
  - High loadings: vehicle\_count, avg\_speed\_kmh (negative), congestion\_level\_encoded
  - Represents traffic volume, congestion, and overall traffic stress.
2. Factor 2 – Atmospheric Weather Severity:
  - High loadings: weather\_severity, humidity, road\_severity
  - Captures weather-related hazards affecting traffic.
3. Factor 3 – Visibility & Flow Quality:
  - High loadings: visibility\_m\_traffic, avg\_speed\_kmh, road\_severity (negative)
  - Indicates quality of traffic flow and visibility conditions.
4. Factor 4 – Seasonal Temperature Patterns:
  - High loadings: temperature\_c, season\_encoded
  - Reflects seasonal variations and temperature effects on traffic.

#### Factor Scores (Gold Layer)

- Factor scores were calculated for each observation using the fitted model.

- Scores were merged with the cleaned dataset to form the Gold Layer, ready for downstream modeling or analysis.

```
First 5 rows of the Gold Layer data (with new factor scores):
   Traffic_Load_Congestion_Score  Atmospheric_Weather_Severity_Score \
0                  -1.032607             -0.906444
1                  -1.722059             -0.735907
2                  -1.253827             -0.862082
3                  -1.186322             -0.857412
4                  -1.032318              0.940379

   Visibility_Flow_Quality_Score  Seasonal_Temperature_Pattern_Score
0                  -0.136459             -0.929442
1                  1.442984              -0.641487
2                  0.140467              -1.094074
3                  0.233843              -1.107967
4                  0.147177              -0.228831
```

### *Conclusion*

- Successfully reduced 15 variables into 4 interpretable latent factors.
- The Gold Layer dataset now contains both original features and factor scores, ready for predictive modeling.
- This approach enables dimensionality reduction, clearer understanding of traffic and weather interactions, and supports urban traffic planning under varying conditions.

## Monte Carlo Simulation

### 1. Data Loading and Method

- Method: Probabilistic data-driven sampling. We draw a few rows per iteration, average their values to get a representative condition, then add random noise. Risk is computed via thresholds.
- Why averaging: It reduces outlier influence; the noise reintroduces variability.

```

df_monte_carlo = pd.read_parquet("merged_data.parquet")
def traffic_jam_risk(temp, rain, humidity, visibility, wind):
    risk = 0

    # Heavy rain
    if rain > 20:
        risk += 0.30

    # Temperature extremes
    if temp < 0 or temp > 35:
        risk += 0.15

    # High humidity
    if humidity > 85:
        risk += 0.15

    # Low visibility
    if visibility < 500:
        risk += 0.30

    # Strong winds
    if wind > 50:
        risk += 0.15

    return min(risk, 1.0)

```

## 2.Traffic jam Risk Function

- Rule-based: Adds fixed weights when hazardous thresholds are crossed (rain, temperature, humidity, visibility, wind).
- Cap at 1.0: Ensures probability stays within [0, 1].
- Effect: Multiple hazards add up; stronger weights (e.g., low visibility) push risk higher.

```

def accident_risk(traffic_jam_prob, visibility, rain):
    risk = 0.1 + traffic_jam_prob * 0.6

    if visibility < 300:
        risk += 0.2

    if rain > 30:
        risk += 0.15

    return min(risk, 1.0)

```

### 3. Accident Risk Function

- Base risk ties accident risk to traffic jam ( $0.1 + 0.6 * \text{traffic jam}$ ).
- Extra penalties: Very low visibility and heavy rain elevate accident likelihood.
- Cap at 1.0: Keeps probability valid.

```
NUM_SIMULATIONS = 10000

simulation_results = []

SAMPLE_SIZE = 5

for i in range(NUM_SIMULATIONS):

    # sample multiple rows instead of one
    sample = df_monte_carlo.sample(SAMPLE_SIZE)
    # Use mean values to reduce sensitivity to a single row
    temp = sample["temperature_c"].mean() + np.random.normal(0, 2)
    rain = max(sample["rain_mm"].mean() + np.random.normal(0, 5), 0)
    humidity = min(max(sample["humidity"].mean() + np.random.normal(0, 5), 0), 100)
    visibility = max(sample["visibility_m_weather"].mean() + np.random.normal(0, 200), 50)
    wind = max(sample["wind_speed_kmh"].mean() + np.random.normal(0, 5), 0)

    traffic_jam_prob = traffic_jam_risk(temp, rain, humidity, visibility, wind)

    accident_prob = accident_risk(traffic_jam_prob, visibility, rain)

    simulation_results.append({
        "simulation_id": i + 1,
        "traffic_jam_probability": traffic_jam_prob,
        "accident_probability": accident_prob
    })
```

### Probabilistic Monte Carlo Loop

- We sample 5 rows, average each variable, and add noise to model uncertainty.
- Threshold rules convert these values to traffic jam probability; accident risk depends on traffic jam, rain, and visibility.
- This method is data-driven; scenarios occur implicitly when thresholds are crossed.

```

Combined_results_df = pd.DataFrame(simulation_results)
Combined_results_df.to_csv("simulation_results.csv", index=False)


plt.hist(Combined_results_df["traffic_jam_probability"], bins=50)
plt.yscale("log")

plt.xlabel("traffic jam Probability")
plt.ylabel("Frequency")
plt.title("Monte Carlo Simulation - traffic jam Probability Distribution")

plt.xticks(
    [0, 0.25, 0.5, 0.75, 1],
    ["No Congestion", "Low", "Medium", "High", "Severe"]
)

plt.grid(True)
plt.show()

```



```


plt.hist(Combined_results_df["accident_probability"], bins=50)
plt.yscale("log")

plt.xlabel("Accident Probability")
plt.ylabel("Frequency")
plt.title("Monte Carlo Simulation - Accident Probability Distribution")

plt.xticks(
    [0, 0.25, 0.5, 0.75, 1],
    ["Low Risk", "Moderate Risk", "High Risk", "Very High Risk", "Severe Risk"]
)

plt.grid(True)
plt.show()

```



### Insights:

Most simulations fall within low to medium traffic jams, indicating that severe traffic jams are uncommon under normal conditions.

Extreme traffic jams and high accident risks are rare but possible, as shown by the long tail in the distributions.


Accident probability increases with higher traffic jams, especially when combined with poor visibility and heavy rainfall.

## Univariate Scenarios for All Hazards

- We run five separate univariate analyses: Heavy Rain, Temperature Extremes, High Humidity, Low Visibility, Strong Winds.

This function simulates a Heavy Rain univariate scenario to analyze traffic risks. It iteratively samples a vehicle count and a rain value (with added Gaussian noise) while keeping other factors neutral. It then calculates the resulting congestion probability and accident probability based on these variables and records the outcomes for report analysis.

```
# Heavy Rain univariate
for i in range(N_UNIV_ALL):
    vehicle_count = sample_traffic()
    rain = float(max(df_monte_carlo["rain_mm"].sample(1).iloc[0] + np.random.normal(0, 3), 0))
    traffic_jam_prob = traffic_jam_risk(NEUTRAL["temp"], rain, NEUTRAL["humidity"], NEUTRAL["visibility"], NEUTRAL["wind"])
    accident_prob = accident_risk(traffic_jam_prob, NEUTRAL["visibility"], rain)
    univ_all_results.append({
```



### Traffic Jam Insights:


The traffic jam probability distribution is right-skewed with mean around 0.25, indicating that heavy rainfall (>20mm) triggers significant congestion risk primarily when combined with high vehicle counts. Most simulations fall in the 0.20-0.30 range, with tail events reaching 0.50+ when multiple hazards combine.

### Accident Risk Insights:

Accident probability shows higher values (mean ~0.35-0.40) compared to traffic jam probability, as it compounds congestion risk with additional rain-specific hazards. The distribution reveals that accident risk escalates significantly in heavy rain scenarios (>30mm), with the model adding 0.15 penalty for extreme precipitation.

This function simulates a **Temperature Extremes univariate scenario** to assess traffic risks. It iteratively **samples a vehicle count** and a **temperature value** (with added Gaussian noise) while keeping all other weather factors neutral. It calculates the resulting **congestion probability** and **accident probability** based on these variables, and records both outcomes for detailed analysis.

```
# Temperature Extremes univariate
for i in range(N_UNIV_ALL):
    vehicle_count = sample_traffic()
    temp = float(df_monte_carlo["temperature_c"].sample(1).iloc[0] + np.random.normal(0, 2))
    traffic_jam_prob = traffic_jam_risk(temp, NEUTRAL["rain"], NEUTRAL["humidity"], NEUTRAL["visibility"], NEUTRAL["wind"])
    accident_prob = accident_risk(traffic_jam_prob, NEUTRAL["visibility"], NEUTRAL["rain"])
```



### Traffic Jam Insights:

Temperature extremes (below 0°C or above 35°C) contribute moderate congestion risk with a weight of 0.15, showing relatively lower impact compared to visibility or rain hazards. The distribution reveals that most temperature-related traffic jam probabilities cluster around 0.20, suggesting that temperature alone rarely causes severe congestion unless paired with high traffic volumes exceeding 3000 vehicles. The narrow distribution indicates consistent, predictable behavior across varied temperature conditions.

### Accident Risk Insights:

Accident probability in temperature scenarios averages around 0.25-0.28, primarily driven by the base congestion risk rather than temperature-specific accident penalties. The distribution shows tight clustering, indicating that temperature extremes affect traffic flow more than directly causing accidents. However, extreme cold (<0°C) can contribute to icy conditions (captured indirectly), while extreme heat (>35°C) may affect driver behavior and vehicle performance.

And so on with the rest of the variables.

## Bonus Dashboard

# Tab 1 data set statistics

localhost:8501

**Dataset Statistics** Monte Carlo Simulation Factor Analysis

## Cleaned Dataset Statistics

### Sample of Cleaned Data :

	traffic_id	date_time	city	area	vehicle_count	avg_speed_kmh	accident_
0	9002	2020-01-01 01:00:00	London	Westminster	267	43.9	
1	9003	2020-01-01 02:00:00	London	Camden	188	59.4	
2	9004	2020-01-01 04:00:00	London	Greenwich	226	47.4	
3	9005	2020-01-01 04:00:00	London	Southwark	340	47.3	
4	9006	2020-01-01 06:00:00	London	Westminster	350	39.8	
5	9007	2020-01-01 06:00:00	London	Islington	787	35.5	
6	9008	2020-01-01 07:00:00	London	Southwark	2253	13.9	
7	9009	2020-01-01 08:00:00	London	Greenwich	3992	19.4	
8	9010	2020-01-01 09:00:00	London	Chelsea	2661	29.5	
9	9011	2020-01-01 11:00:00	London	Kensington	1815	36.3	

### Shape of dataset:

(4162, 19)

### Key Metrics

Total Records	Total Accidents	Average Speed (km/h)
4162	1465	32.5

### Summary Statistics:

	traffic_id	vehicle_count	avg_speed_kmh	accident_count	visibility_m_traffic	weather_id	te
count	4162	4162	4162	4162	4162	4162	4162
mean	11394.5425	1426.4786	32.5017	0.352	7394.4993	7395.0387	

Streamlit

localhost:8501

### Shape of dataset:

(4162, 19)

### Key Metrics

Total Records	Total Accidents	Average Speed (km/h)
4162	1465	32.5

### Summary Statistics:

	traffic_id	vehicle_count	avg_speed_kmh	accident_count	visibility_m_traffic	weather_id	te
count	4162	4162	4162	4162	4162	4162	4162
mean	11394.5425	1426.4786	32.5017	0.352	7394.4993	7395.0387	
std	1441.0792	1095.8349	11.8521	2.4547	3422.1541	1441.0831	
min	9002	45	0	0	0	5002	
25%	10137.75	405	24.4	0	4775.75	6138.75	
50%	11404.5	1256.5	32.7	0	9917	7405	
75%	12641.75	1930	41.6	0	10006	8642.75	
max	13899	4332	64.8	55	10999	9899	

### Missing Values:

Column Name	Missing Values
0 traffic_id	0
1 date_time	0
2 city	0
3 area	0
4 vehicle_count	0
5 avg_speed_kmh	0
6 accident_count	0
7 congestion_level	0
8 road_condition	0

max	13899	4332	64.8	55	10099	9899
-----	-------	------	------	----	-------	------


### Missing Values:

Column Name	Missing Values
0 traffic_id	0
1 date_time	0
2 city	0
3 area	0
4 vehicle_count	0
5 avg_speed_kmh	0
6 accident_count	0
7 congestion_level	0
8 road_condition	0
9 visibility_m_traffic	0

Select a numeric column

traffic\_id


Distribution of traffic\_id



Select a numeric column

traffic\_id

Distribution of traffic\_id




## Distribution with Summary

Select numeric column

traffic\_id

Distribution of traffic\_id





Statistics

	traffic_id
count	4162
mean	11394.5425
std	1441.0792
min	9002
25%	10137.75
50%	11404.5
75%	12641.75
max	13899



## Distribution with Summary





### Distribution with Summary

Select numeric column

traffic\_id

Distribution of traffic\_id

Statistics


	traffic_id
count	4162
mean	11394.5425
std	1441.0792
min	9002
25%	10137.75
50%	11404.5
75%	12641.75
max	13899

- localhost:8501

Select a numeric column

accident\_count

Distribution of accident\_count




### Distribution with Summary

Select numeric column


temperature\_c

Distribution of temperature\_c


Statistics



Tab 2 Monte carlo:



## Tab 3 Factor Analysis:



## Top & Bottom Records

Number of records to display

5

20

Highest Factor Scores

	Visibility and Smooth Traffic Flow
3911	3.4211
4060	2.8706
446	2.8107
1838	2.7719
3532	2.7298

Lowest Factor Scores

	Visibility and Smooth Traffic Flow
2238	-2.9451
3364	-2.8765
2864	-2.7848
642	-2.7229
2311	-2.6856

1838 3532	2.7719 2.7298	642 2311	-2.7229 -2.6856
--------------	------------------	-------------	--------------------

## Factor Interpretation (Loadings)

Select Factor for Interpretation


Factor 1

Top contributing variables

5

Variable	Factor 1	Absolute Loading
0 vehicle_count	0.923	0.923
1 avg_speed_kmh	-0.792	0.792
11 congestion_level_encoded	0.708	0.708
10 hour	0.392	0.392
4 temperature_c	0.146	0.146

### Top Variables Contributing to Factor 1



## Factor Interpretation (Loadings)

Select Factor for Interpretation

Factor 3

Top contributing variables

3

Variable	Factor 3	Absolute Loading
3 visibility_m_traffic	0.634	0.634
1 avg_speed_kmh	0.566	0.566
13 road_severity	-0.486	0.486
11 congestion_level_encoded	-0.433	0.433
0 vehicle_count	0.228	0.228
8 visibility_m_weather	0.219	0.219
10 hour	0.081	0.081
4 temperature_c	0.07	0.07
5 humidity	-0.07	0.07
2 accident_count	-0.065	0.065

### Top Variables Contributing to Factor 3

